



FPGA-based Distributed Edge Training of SVM

Jyotikrishna Dass, Yashwardhan Narawane, Rabi Mahapatra, Vivek Sarin
Texas A&M University



COMPUTER SCIENCE
& ENGINEERING
TEXAS A&M UNIVERSITY

Overview

Goal

- ✓ To design and implement distributed training of machine learning, here, Support Vector Machines (SVM), on multiple FPGA system
- ✓ To reduce network communication while achieving fast training, memory-efficiency, and energy savings

Motivation

- SVM training is computationally expensive with high memory requirement for *kernel* matrix
- Traditional SVM training accelerators are based on inherently sequential algorithms using a single FPGA board
- Need to distribute and accelerate training on edge where the data is generated and stored across multiple devices

Background

Support Vector Machines

- Supervised machine learning algorithm for classification and regression problems
- Mathematically, a quadratic programming problem which solves for maximal separating hyperplane as a classifier

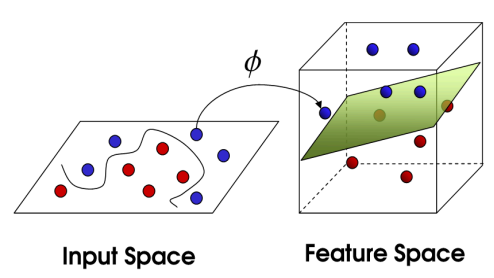


Fig. Kernel SVM: Learning hyperplane in higher feature dimension

training dataset, $\mathcal{D} = \{(x_i, y_i), i = 1, \dots, n\}$
input data matrix, $X = \{x_i \in \mathbb{R}^d, i = 1 \dots n\}$, d -dimensional space
class label vector, $y = \{y_i \in \{-1, 1\}, i = 1 \dots n\}$

SVM dual

$$\min_{\alpha} \frac{1}{2} \alpha^T (\text{diag}(y) \times K \times \text{diag}(y)^T) \alpha + \frac{1}{2} \alpha^T \left(\frac{1}{2C} l_n \right) \alpha + e^T \alpha$$

subject to $-\alpha_i \leq 0_n$

where, α is a vector of dual variables

$e = -1_n$

$C > 0$ is penalty parameter for misclassification

$K = \{k(x_i, x_j), \forall i, j = 1 \dots n\}$ is positive definite matrix (mostly $k()$ represents the Mercer kernel function - linear/non-linear)

Substitute, $K \approx AA^T$, $A \in \mathbb{R}^{n \times k}$ and $k \ll n$.

Define, $\hat{A} = \text{diag}(y) \times A$

Substitute, $\hat{A} = QR$, where, $Q \in \mathbb{R}^{n \times n}$ is Orthogonal matrix and $R \in \mathbb{R}^{n \times k}$ is Upper Triangular matrix

QRSVM

$$\min_{\hat{\alpha}} \frac{1}{2} \hat{\alpha}^T (RR^T + \frac{1}{2C} l_n) \hat{\alpha} + (\hat{e})^T \hat{\alpha}$$

subject to $-\hat{\alpha}_i \leq 0_n$

where, $\hat{\alpha} = Q^T \alpha$, $\hat{e} = Q^T e$

Define, $F = -(R_g R_g^T + \frac{1}{2C} l_n)$ and $\hat{n} = \hat{e}$

Step 1: Minimization of Lagrangian - In Parallel

At edge unit, i

$$\hat{\alpha}_i^{k+1} = F_i^{-1} (-\hat{n}_i^k + \hat{\theta}_i)$$

where,

$$F_i^{-1} = \begin{cases} F_1^{-1} & \text{if } i = 1 \\ -2C & \text{if } i = 2..p \end{cases}$$

Step 2: Dual variable update - In Parallel

At edge unit, i

$$\hat{\beta}_i^{k+1} = \hat{\beta}_i^k + \eta^* (-\hat{\alpha}_i^{k+1})$$

η^* is the Optimal step size

$$\hat{\beta}^k = Q^T \beta^k$$

Algorithmic Design

Distributed QR Decomposition

Algorithm 1 Distributed QR decomposition

- 1: $k \leftarrow \text{rank}$
- 2: **for** each edge i **do**
- 3: $\hat{A}_i \leftarrow$ local partitioned data
- 4: Parallel Compute $\{q_i\}, R_i \leftarrow \hat{A}_i$ ➤ Algorithm 2
- 5: GATHER $(R_i)_{k \times k}$ at Master unit
- 6: **end for**
- 7: $\hat{A}_g \leftarrow$ gathered or stacked $(R_i)_{k \times k}$
- 8: Compute $\{q_f\}, R_f \leftarrow \hat{A}_g$ at Master unit ➤ Algorithm 2
- 9: Use $(R_f)_{k \times k}$

Algorithm 2 $\{q_i\}, R_i \leftarrow \hat{A}_i$, via Householder algorithm

- 1: $Q_{\hat{A} \times k}, (\hat{A}_i)_{\hat{A} \times k}$ ➤ \hat{n} : samples per edge
- 2: **for** $j \leftarrow 1$ **to** k **do**
- 3: $q_{ij} \leftarrow \hat{A}_i(j : \hat{n}, j)$
- 4: $q_{ij}(1) \leftarrow q_{ij}(1) + \text{sign}(q_{ij}(1)) \times \|q_{ij}\|_2$ ➤ scalar update
- 5: $q_{ij} \leftarrow \frac{q_{ij}}{\|q_{ij}\|_2}$ ➤ vector normalization
- 6: $\hat{A}_i(j : \hat{n}, j : k) \leftarrow \hat{A}_i(j : \hat{n}, j : k) - 2q_{ij} < q_{ij}, \hat{A}_i(j : \hat{n}, j : k) >$ ➤ Algorithm 4
- 7: $R_i = \hat{A}_i(j : \hat{n}, j : k)$
- 8: **end for**
- 9: $\{q_i\} \leftarrow [q_{i1}, q_{i2}, \dots, q_{ik}]$ ➤ set of k -reflectors




Fig. Computational flow graph for distributed QR decomposition

Algorithm 4 Computing Step 6 in Algorithm 2

- 1: $\hat{A}_i(j : \hat{n}, j : k) \leftarrow \hat{A}_i(j : \hat{n}, j : k) - 2q_{ij} < q_{ij}, \hat{A}_i(j : \hat{n}, j : k) >$
- 2: **for** $m \leftarrow j$ **to** k **do**
- 3: $a_{BRAM} \leftarrow \hat{A}_i(j : \hat{n}, m)$ ➤ Load into BRAM
- 4: Compute $\text{sum} \leftarrow q_{ij} \cdot a_{BRAM}$ ➤ Inner Product
- 5: $a_{BRAM} \leftarrow a_{BRAM} - 2 \times \text{sum} \times q_{ij}$ ➤ saxpy
- 6: $\hat{A}_i(j : \hat{n}, m) \leftarrow a_{BRAM}$ ➤ Write to DDR
- 7: **end for**

Parallel Dual Ascent

Algorithm 3 Parallel Dual Ascent

- 1: iteration $t \leftarrow 0$
- 2: **for** each edge i **do**
- 3: **while** error $>$ threshold **do**
- 4: Parallel Compute: $\hat{\alpha}_i^{t+1}$ ➤ alpha update
- 5: Parallel Compute: $\hat{\beta}_i^{t+1}$ ➤ dual update
- 6: Compute $\beta_i \leftarrow Q \hat{\beta}_i$ ➤ $\beta_{[ij]} \leftarrow \beta_{[ij]} - 2q_{ij} < q_{ij}, \beta_{[ij]} >$ $j = k$ to 1
- 7: Parallel Compute $\beta_i \leftarrow \max\{0, \beta_i\}$ ➤ Non-negativity check
- 8: Compute $\beta_i \leftarrow Q^T \beta_i$ ➤ $\beta_{[ij]} \leftarrow \beta_{[ij]} - 2q_{ij} < q_{ij}, \beta_{[ij]} >$ $j = 1$ to k
- 9: error $\leftarrow \|\hat{\beta}_i^{t+1} - \hat{\beta}_i^t\|_1$
- 10: $t \leftarrow t + 1$
- 11: **end while**
- 12: **end for**

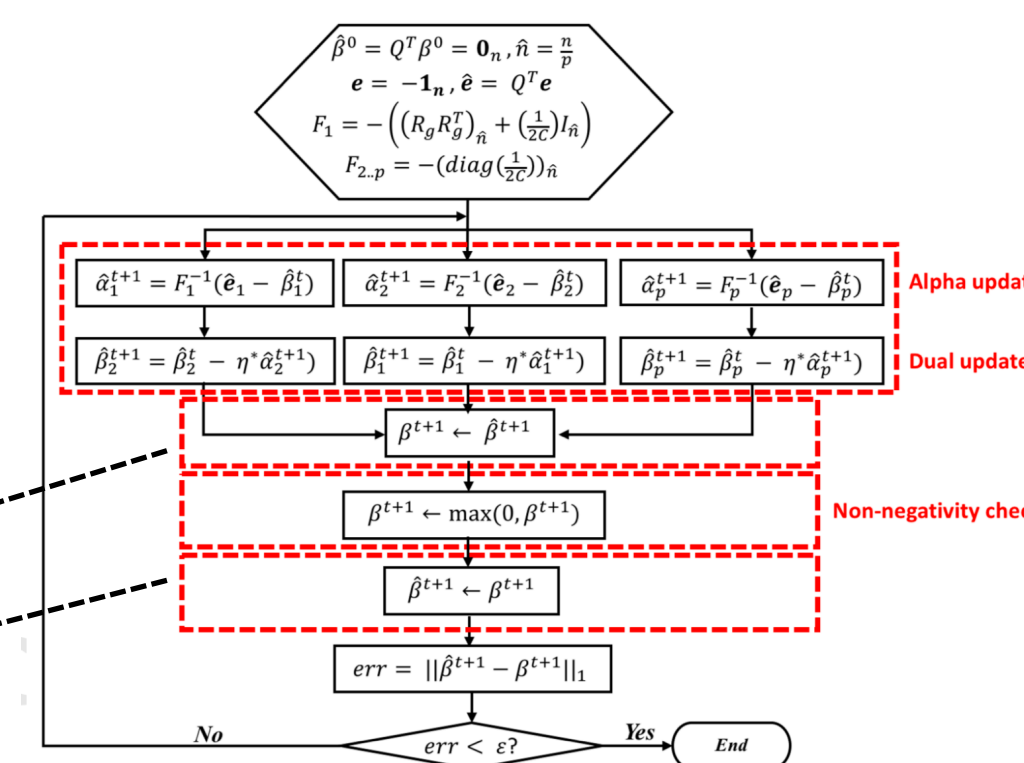


Fig. Computational flow graph for parallel dual ascent

Hardware Implementation

Basic Computational Kernels

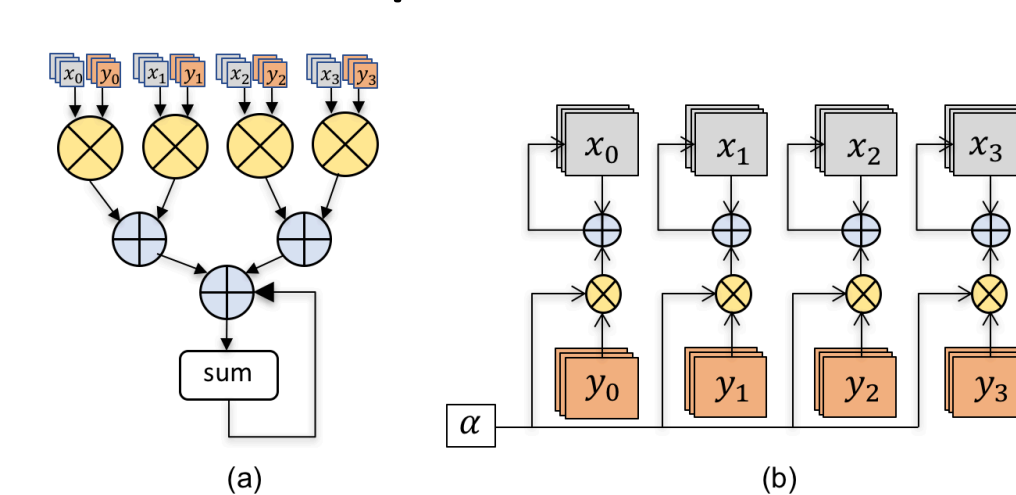


Fig. SIMD design for computational kernel (a) Inner Product $\text{sum} = \langle x, y \rangle$ via binary reduction tree (b) Scaled vector addition SAXPY, $(x = x + \alpha y)$ via fine-grained parallelism. The vectorized kernels operate on $W = \text{floor}(N/B)$ data samples in each pass

Data Layout + Memory Interface




Fig. Data layout in column-major order. Memory-IP interface for on-chip Block RAM (full duplex) is called *bram* and that for off-chip DDR (half-duplex) is called *AXI Master*. $N_{bram} = 1024$ bits, $N_{axi} = 2048$ bits

Increasing Throughput




Fig. Doubling the throughput for Inner Product $\langle x, y \rangle = \langle x_1, y_1 \rangle + \langle x_2, y_2 \rangle$

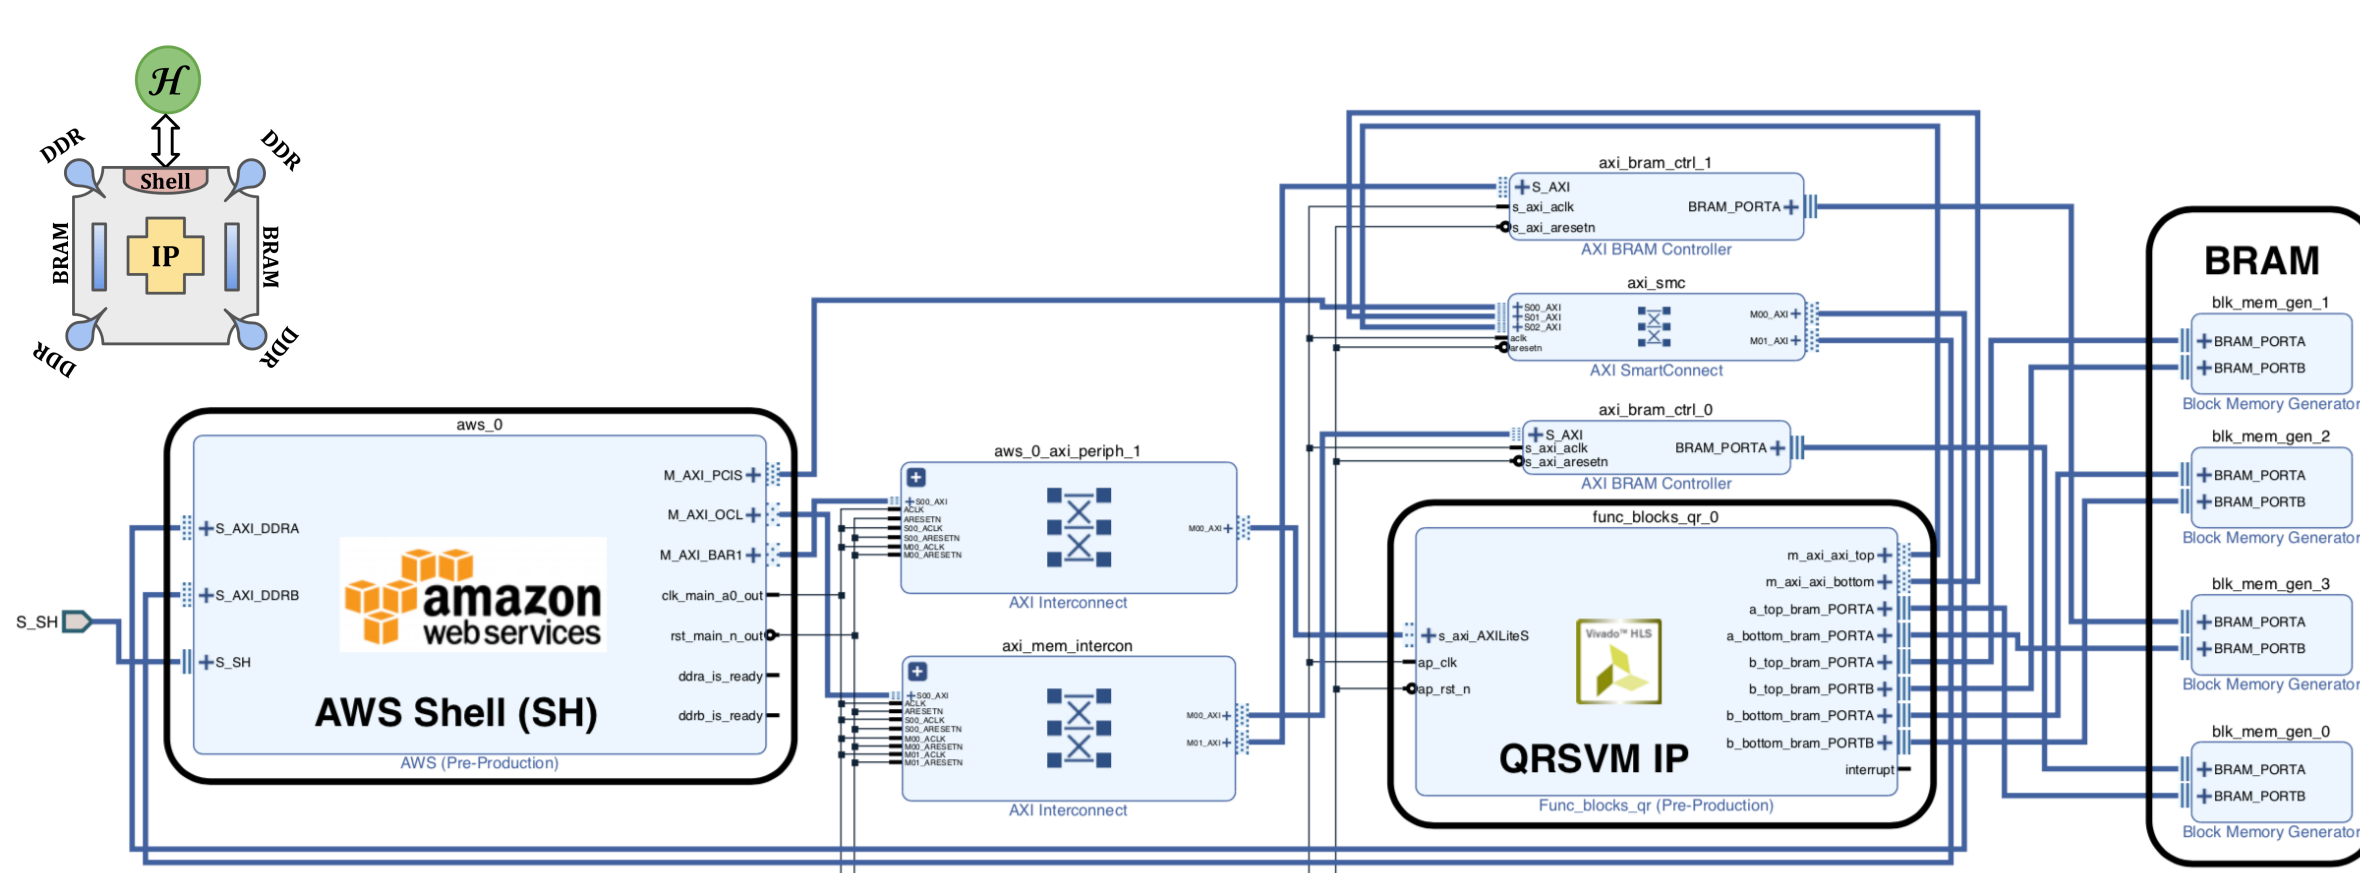
Memory Assignment

Kernel	BRAM	DDR
Inner Product	sum	x, y
SAXPY	x	y

FPGA Xilinx Virtex xcvu9p-flgb2104-2-i


Resource	BRAM	DSP	FF	LUT
Available	2160	6840	2.36M	1.18M
Used	1405	1221	545K	450K
%	65%	18%	23%	38%

FPGA-Architecture for distributed SVM training



Multiple FPGA Network

- Each edge device comprises of FPGA IP logic + Host processor
- Illustration for $p = 4$ edges in a network
- A single QRSVM IP is synthesized per FPGA device to operate at clock frequency of 125 MHz with 39 Watts of power
- Computational workload is entirely with the FPGA while communication is through Host processor (PCIe)
- Each FPGA handles maximum of 256K samples



Experimental Results and Discussions

Hardware Platform

Proof-of-Concept implemented on Amazon AWS F1 instance with $p=\{1, 2, 4, 8\}$ 16nm Xilinx Virtex Ultrascale+ VU9P FPGA units forming a multiple FPGA network

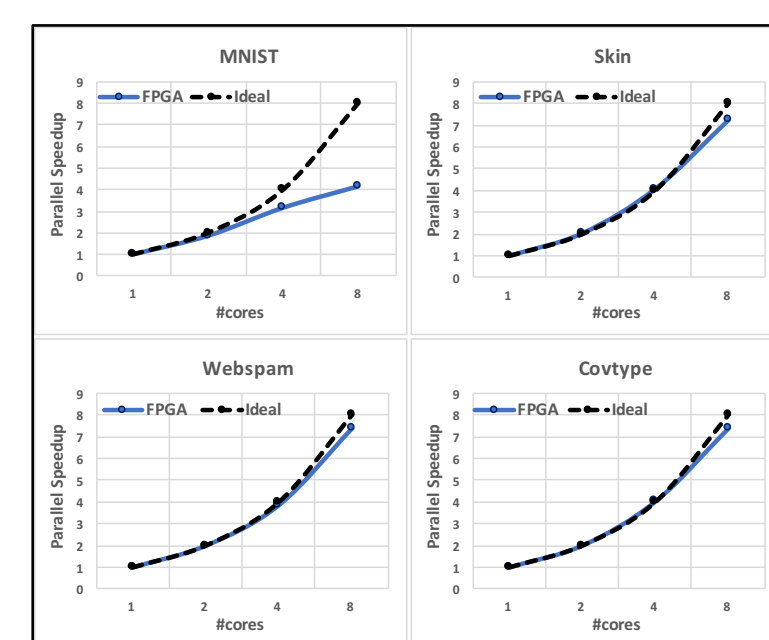
Dataset Description

Benchmark	Application	#samples (n)	#features (d)	k-rank
MNIST	Image	60,000	780	128
Skin	Health	200,000	3	64
Webpsam	Email	350,000	254	128
Covtype	Geography	464,810	54	64
SUSY	Physics	2,000,000	18	128

Training Time

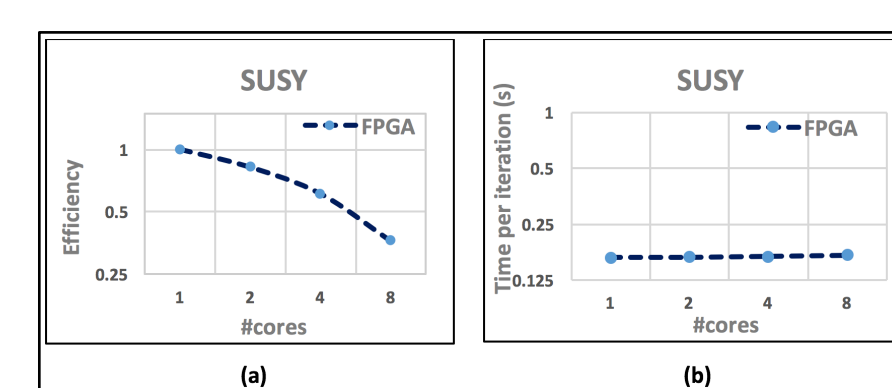
#units	#iterations	MNIST ($C = 1, \gamma = 2^{-6}, \eta^* = 0.9$)					
p	t	T_{QR}	T_{DA}	T_{FPGA}^p	comp	comm	
1	181	3.42	7.49	10.92	99.9%	0.1%	
2	181	1.76	4.07	5.84	99.8%	0.2%	
4	182	0.93	2.51	3.58	96%	4%	
8	182	0.46	2.14	2.61	99.6%	0.4%	
#units	#iterations	Skin ($C = 1, \gamma = 2^{-8}, \eta^* = 0.9$)					
p	t	T_{QR}	T_{DA}	T_{FPGA}^p	comp	comm	
1	67441	2.80	4533	4536	99.9%	0.1%	
2	64424	1.46	2226	2228	99.9%	0.1%	
4	59761	0.74	1107	1108	99.9%	0.1%	
8	54744	0.38	625	626	99.9%	0.1%	
#units	#iterations	Webpsam ($C = 1, \gamma = 1, \eta^* = 0.9$)					
p	t	T_{QR}	T_{DA}	T_{FPGA}^p	comp	comm	
1	-	-	-	-	-	-	
2	566	9.80	66.20	76.14	99.8%	0.2%	
4	564	4.88	34.40	39.36	99.8%	0.2%	
8	569	2.60	17.92	20.59	99.7%	0.3%	
#units	#iterations	Covtype ($C = 1, \gamma = 2^3, \eta^* = 0.9$)					
p	t	T_{QR}	T_{DA}	T_{FPGA}^p	comp	comm	
1	-	-	-	-	-	-	
2	1125	3.35	88.02	91.45	99.9%	0.1%	
4	1080	1.70	43.58	45.36	99.8%	0.2%	
8	1068	0.80	23.80	24.75	99.4%	0.6%	
#units	#samples	SUSY ($C = 1, \gamma = 2^{-8}, \eta^* = 0.9$)					
p	n	T_{QR}	T_{DA}	T_{FPGA}^p	comp	comm	
1	250K	14.01	94.04	108.08	99.9%	0.1%	
2	500K	14.04	116.8	131.02	99.8%	0.2%	
4	1M	14.07	162.1	176.18	99.9%	0.1%	
8	2M	14.14	285.47	299.63	99.9%	0.1%	

Strong Scaling Analysis



- Achieves near **linear parallel speedup** for larger datasets Skin, Webpsam, Covtype
- For small dataset MNIST, going beyond $p = 4$ seems to be overkill

Weak Scaling Analysis



- **Parallel Dual Ascent** is computationally dominant than Distributed QR decomposition
- Near negligible communication overhead **<1%**
- Workload per FPGA fixed at 250K samples
- (a) T_{QR} is constant while $T_{DA} \uparrow$ with #iterations
- (b) (T_{FPGA}/t) is constant as desired

Energy Analysis

- Under strong scaling, the proposed FPGA design follows the **ideal energy consumption** trend that is **constant** across #FPGA units.
- Validates **fully parallel** implementation

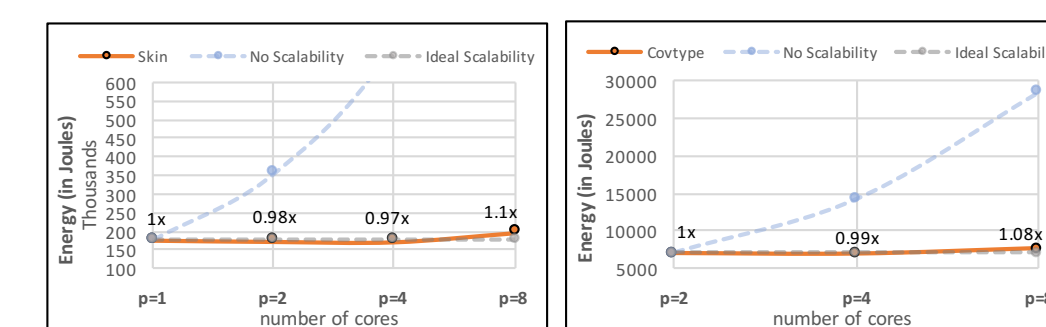


Fig. Energy consumption under Strong Scaling Skin and Covtype

- Under weak scaling, the ideal energy consumption trend is linear while no scalability trend is quadratic.
- The proposed design is closer to being linear than quadratic. Aberration at $p=8$ due to large #iterations for fine tuning model with increasing overall problem size
- (Energy/p) is nearly **constant** as expected with uniform workload per device




Fig. (a) Total Energy (b) Energy per core consumption under Weak Scaling for SUSY

Future Work

- ❑ To design and implement distributed training for Deep Learning models for resource constrained devices with limited memory and low power
- ❑ To explore online/incremental learning capabilities for machine learning models at edge

References

1. Burges, Christopher JC. "A tutorial on support vector machines for pattern recognition." Data mining and knowledge discovery 2, no. 2 (1998): 121-167.
2. Papadonikolakis, Markos, and Christos-Savvas Bouganis. "A scalable FPGA architecture for non-linear SVM training." In 2008 International Conference on Field-Programmable Technology, pp. 337-340. IEEE, 2008.
3. Cadambi, Srihari, Igor Durdanovic, Venkata Jakkula, Murugan Sankaradass, Eric Cosatto, Srmat Chakradhar, and Hans Peter Graf. "A massively parallel FPGA-based coprocessor for support vector machines." In 2009 17th IEEE Symposium on Field Programmable Custom Computing Machines, pp. 115-122. IEEE, 2009.
4. Rabieah, Mudhar Bin, and Christos-Savvas Bouganis. "FPGA based nonlinear support vector machine training using an ensemble learning." In 2015 25th International Conference on Field Programmable Logic and Applications (FPL), pp. 1-4. IEEE, 2015.
5. Si, Si, Cho-Jui Hsieh, and Indejit S. Dhillon. "Memory efficient kernel approximation." The Journal of Machine Learning Research 18, no. 1 (2017): 682-713.
6. Dass, Jyotikrishna, Vivek Sarin, and Rabi N. Mahapatra. "Fast and Communication-Efficient Algorithm for Distributed Support Vector Machine Training." IEEE Transactions on Parallel and Distributed Systems (2018).